
Machine Problem 0 - xv6 Setup

CSIE3310 - Operating Systems
National Taiwan University

Total Points: 100
Release Date: February 20
Due Date: February 27, 23:59:59
TA hours: Tue. & Thu. 13:00-15:00 before the due date, CSIE Building R401

Contents

1	Discussion Policy	1
2	Summary	1
3	Docker Installation	2
3.1	Why Docker?	2
3.2	Supported Platforms	2
3.3	Installation Testing	2
4	Launching xv6	3
4.1	Launching the Docker Image of MP0	3
4.2	QEMU Introduction	3
4.3	Launching QEMU	4
5	Example MP	4
5.1	tree command and mp0 command	5
5.2	mp0 Executable	5
5.3	Detailed Steps	5
5.4	Testcase Specifications	5
5.5	Sample Outputs	6
5.5.1	Sample Testcase 1	6
5.5.2	Sample Testcase 2	6
5.5.3	Sample Testcase 3	6
5.6	Public Testcases	7
5.7	Grading Procedure	7
6	Submission	7
6.1	Folder Structure after Unzipping	7
6.2	Grading Policy	8
7	References	8

1 Discussion Policy

If you have any question about this machine problem, please post it on the corresponding NTUCOOL discussion. We have opened a discussion dedicated to MP0. However, if you have some special requests, you can send an email to ntuos@googlegroups.com. Note that we might not reply if you ask about a general problem by sending email.

2 Summary

xv6 is an example kernel created by MIT for pedagogical purposes. We will study xv6 to get familiar with the main concepts of operating systems. The reference book of xv6 is [xv6: a simple, Unix-like teaching operating system](#). You will learn to set up the environment for xv6 and develop a custom mp0 command in this MP.

3 Docker Installation

If you have installed Docker in a Unix environment, you can directly go to the next step. If your operating system is Windows, we strongly suggest you to [install WSL2](#) and [run Docker in WSL2](#).

3.1 Why Docker?

To run an arbitrary OS on your machine, you need virtualization. Virtualization is the process of running a virtual instance of a computer system in a layer abstracted from the actual hardware. A virtual machine (VM) is the emulated equivalent of a computer system that runs on top of another system. However, VM runs a guest OS with virtual access to host resources through a hypervisor. It generally incurs lots of overhead beyond what is being consumed by your application logic. By contrast, the container runs a discrete process, taking no more memory than any other executable, making it lightweight. Docker is a platform for you to build and run with containers. We also leverage the advantage of virtualization to standardize the environment of your homework, making the problems independent from both architectures and operating systems of your machines.

3.2 Supported Platforms

Docker is available on Windows, Mac, and various Linux platforms. You can find your preferred operating system below and follow the instructions of the installation guide.

- [Docker Desktop for Windows](#)
- [Docker Desktop for Mac](#)
- [Docker Engine for Ubuntu](#)
- [Docker Engine for Debian](#)
- [Docker Engine for CentOS](#)
- [Docker Engine for Fedora](#)

We suggest you install Docker on Linux platforms because the Docker commands we provide in MPs are based on Linux platforms. We do not guarantee to answer Docker issues on other platforms.

3.3 Installation Testing

To check whether you have installed Docker successfully, run the hello-world image.

```
$ docker run hello-world
```

You should see something like this

```
Hello from Docker!
This message shows that your installation appears to be working correctly.

To generate this message, Docker took the following steps
1. The Docker client contacted the Docker daemon.
2. The Docker daemon pulled the "hello-world" image from the Docker Hub.
   (arm64v8)
3. The Docker daemon created a new container from that image which runs the
   executable that produces the output you are currently reading.
4. The Docker daemon streamed that output to the Docker client, which sent it
   to your terminal.

To try something more ambitious, you can run an Ubuntu container with
$ docker run -it ubuntu bash

Share images, automate workflows, and more with a free Docker ID
https://hub.docker.com/

For more examples and ideas, visit:
https://docs.docker.com/get-started/
```

4 Launching xv6

4.1 Launching the Docker Image of MP0

1. Download the MP0.zip from NTUCOOL, unzip it, and enter it.

```
$ unzip MP0.zip  
$ cd mp0
```

2. Pull the Docker image from Docker Hub

```
$ docker pull ntuos/mp0
```

Note that the image only supports the following CPU architectures

- x86_64 or amd64
- arm64

You can check the architecture of your CPU by running the following command

```
$ arch
```

If you are not able to use the image we provided, please send an email to the TA and let us know the output of your arch command.

3. Use docker run to start the process in a container and allocate a TTY for the container process

```
$ docker run -it -v $(pwd)/xv6:/home/os_mp0/xv6 ntuos/mp0
```

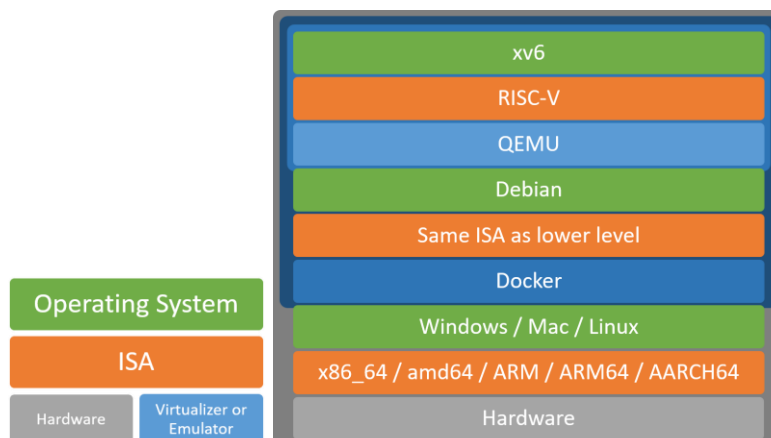
This command will make the Docker container interactive and pretend to be a shell. You may notice that we mount a volume with -v. Volumes are often a better choice than persisting data in a container's writable layer because a volume does not increase the size of the containers using it, and the volume's contents exist outside the life-cycle of a given container. We adopt it here to give you the flexibility to do coding outside the container with your favorite editor rather than coding with vim inside. Of course, you can use vim if you really love it.

4. Please check the environment in the Docker container

```
$ cat /etc/os-release
```

4.2 QEMU Introduction

xv6 is implemented in ANSI C for a multi-core RISC-V system, but most of our machines are not RISC-V systems. Therefore, we need QEMU to help us launch xv6 on non-RSIC-V architectures. QEMU is an emulator and virtualizer that can perform hardware virtualization. An instruction set architecture (ISA) is an abstract model of a computer, also referred to as computer architecture. The ISA serves as the interface between software and hardware.



4.3 Launching QEMU

If you successfully run the Docker container, you should be in the `xv6` directory. In the `xv6` directory, build and run `xv6` on QEMU by running `make qemu`. The behavior of `make qemu` is defined by the Makefile we provided for you.

```
$ make qemu
...

xv6 kernel is booting

hart 2 starting
hart 1 starting
init: starting sh
$
```

If you type `ls` at the prompt, you should see an output similar to the following

```
$ ls
-          1 1 1024
..         1 1 1024
README    2 2 2059
cat       2 3 24096
echo      2 4 22920
grep      2 5 27400
init      2 6 23656
kill      2 7 22880
ln        2 8 22720
ls        2 9 26288
mkdir     2 10 23016
rm        2 11 23008
sh        2 12 41832
stressfs  2 13 23880
grind     2 14 37984
wc        2 15 25208
zombie    2 16 22256
testgen   2 17 25664
console   3 18 0
```

These are the files that `mkfs` includes in the initial file system; most are executable programs. Now you have successfully set up `xv6`. You are welcome to play around in preparation for future MPs. Here are some useful commands for you

- `Ctrl-p`: Print information about each process. If you try it now, you will see two lines: one for `init` and the other for `sh`.
- `Ctrl-a x`: Quit QEMU. Note that you should hold `Ctrl` and press `a`, release both `Ctrl` and `a`, and then press `x` to quit.

5 Example MP

Below we provide an easy example MP. If you find it hard, you should think twice before taking this course or work harder. We assume

- You are good at coding with C.
- You are familiar with System Programming.
- You are willing to learn the competency of looking up reference books and tracing source code.

The ability to search the solution on Stack Overflow will not make you escape from this. If you feel disinclined to learn them, you may also feel terrible while doing MPs.

5.1 tree command and mp0 command

`tree` is a command used to list the contents of directories in a tree-like format. In this MP, we will implement a command called `mp0` to mimic the traversing behavior of the `tree` command. However, we do not need to print the traversed file in a tree-like format. Instead, you are asked to count the number of occurrences of a given key in each traversed path.

5.2 mp0 Executable

To develop a custom xv6 command, you need to modify `mp0/xv6/Makefile` so that the system will build the executable of `mp0` after you run `make qemu`.

5.3 Detailed Steps

1. The process `mp0` reads two argument from the command line.

```
$ mp0 <root_directory> <key>
```

2. The process forks a child.
3. The child process has to traverse the files and directories under `<root_directory>`, output the path and the number of occurrences of the given key, and analyze the total number of traversed files and directories. The testcases only contain files and directories (no devices). The path and the number of occurrences should be separated by a space. The traversing order should be the same as the `ls` command. Note that if the path is `aa/a`, the occurrence should be 3.

```
<path> <occurrence>
```

Note that you should separate the `<root_directory>` and its sub-directory with a single `"/`. That is, if `<root_directory>` contains one or more `"/`, then you should keep all, append a `"/`, and then append the name of the sub-directory. Sample outputs are provided in Section 5.5.

However, if the `<root_directory>` does not exist or is not a directory, you should print the following line:

```
<root_directory> [error opening dir]
```

4. The child process sends two integers, `file_num` and `dir_num`, to the parent process through a pipe.

- `file_num`: the number of traversed files
- `dir_num`: the number of traversed directories

5. The parent process reads the integers from the pipe and prints

```
<dir_num> directories, <file_num> files
```

`<file_num>` is the number of traversed files, and `<dir_num>` is the number of traversed directories. The `<root_directory>` should not be count. You should separate the child process and the parent process outputs with a blank line. You can print the blank line in either the child process or the parent process. However, if you do not use a pipe to receive `dir_num` and `file_num` and print out the result in the parent process, you will be regarded as CHEATING and get 0 points in this MP.

5.4 Testcase Specifications

- `<key>` will only contain a lowercase alphabet (a-z).
- We define the depth of the `<root_directory>` as 0, files and directories in the `<root_directory>` as 1. The depth of the traversed files and directories will be smaller than 5. For example, we may traverse the file `a/b/c/d/e`. But the testcases will not contain the file `a/b/c/d/e/f`.
- The total number of files and directories will not be larger than 20.
- The maximum length of each directory name and file name is 10.

- You only have to consider files and directories. We will not generate a testcase that contains a device.
- The testcases are guaranteed to have sufficient parameters when executing the command. We will not execute the mp0 command without providing <root directory> and <key>.

5.5 Sample Outputs

5.5.1 Sample Testcase 1

Assume that there is a directory os2023 with the following structure

- three sub-directories: d1, d2, and d3.
- d2 has two sub-directories a, b, and one file c.
- d3 has one sub-directory a and one file b

You can generate the structure by running the command testgen. Also, we assume that the given key is “d”. The output of the command mp0 os2023 d should be as follows:

```
$ testgen
$ mp0 os2023 d
os2023 0
os2023/d1 1
os2023/d2 1
os2023/d2/a 1
os2023/d2/b 1
os2023/d2/c 1
os2023/d3 1
os2023/d3/a 1
os2023/d3/b 1

6 directories, 2 files
```

5.5.2 Sample Testcase 2

The output of the command mp0 os2023/ d should be as follows:

```
$ testgen
$ mp0 os2023/ d
os2023/ 0
os2023//d1 1
os2023//d2 1
os2023//d2/a 1
os2023//d2/b 1
os2023//d2/c 1
os2023//d3 1
os2023//d3/a 1
os2023//d3/b 1

6 directories, 2 files
```

5.5.3 Sample Testcase 3

Assume that os2202 does not exist

```
$ mp0 os2202 d
os2202 [error opening dir]

0 directories, 0 files
```

5.6 Public Testcases

You can get 40 points (100 points in total) if you pass all public testcases. You can judge the code by running the following command in the docker container (not in xv6; this should run in the same place as `make qemu`). Note that you should only modify `xv6/Makefile` and `xv6/user/mp0.c`. We do not guarantee that you can get public points if you modify other files to pass all testcases.

```
$ make grade
```

If you successfully pass all the public testcases, the output should be the same as below

```
== Test mp0 command with public testcase 0 (5%) ==
mp0 command with public testcase 0: OK (10.9s)
== Test mp0 command with public testcase 1 (5%) ==
mp0 command with public testcase 1: OK (0.7s)
== Test mp0 command with public testcase 2 (5%) ==
mp0 command with public testcase 2: OK (0.7s)
== Test mp0 command with public testcase 3 (5%) ==
mp0 command with public testcase 3: OK (0.9s)
== Test mp0 command with public testcase 4 (10%) ==
mp0 command with public testcase 4: OK (1.1s)
== Test mp0 command with public testcase 5 (10%) ==
mp0 command with public testcase 5: OK (0.9s)
Score: 40/40
```

If you want to know the details about the testcases, please check `xv6/grade-mp0` and `xv6/user/testgen.c`

Hints:

1. Remember to close the file descriptor.
2. Check `xv6/user/grind.c` and `xv6/user/lsc.c` for function usage.
3. Set the buffer size to the maximum possible length of a path when writing a recursive function. Otherwise, you might encounter some problems such as function stack overflow.

5.7 Grading Procedure

We will grade your submission by running `make grade` (but with additional private testcases).

6 Submission

Please compress your xv6 source code as `<whatever>.zip` and upload it to NTUCOOL. The filename does not matter since NTUCOOL will rename your submissions. Never compress files we do not request, such as `.o`, `.d`, `.asm` files. You can run `make clean` in the container before you compress. Make sure your xv6 can be compiled by `make qemu`.

6.1 Folder Structure after Unzipping

We will unzip the submission by running `unzip *.zip`. The structure AFTER UNZIPPING should be

```
<student_id>
|
+-- mp0.c
|
+-- Makefile
```

Note that the English letters in the `<student id>` must be lowercase. For example, it should be `d08922025` instead of `D08922025`. When grading your submission, we will provide other files and copy your `Makefile` and `mp0.c` to `mp0/xv6/` and `mp0/xv6/user/`, respectively.

6.2 Grading Policy

- You will get 0 points if we cannot compile your submission.
- You will be deducted 10 points if the folder structure is wrong. Using uppercase in the <student id> is also a type of wrong folder structure.
- If your submission is late for n days, your score will be $\max(\text{raw_score} - 20 \times \lceil n \rceil, 0)$. Note that you will not get any points if $\lceil n \rceil \geq 5$.

7 References

- [1] xv6, a simple Unix-like teaching operating system. <https://pdos.csail.mit.edu/6.828/2012/xv6.html>
- [2] Docker: Empowering App Development for Developers. <https://docs.docker.com/>
- [3] RISC-V: The Free and Open RISC Instruction Set Architecture. <https://riscv.org/>
- [4] QEMU, the FAST! processor emulator. <https://www.qemu.org/>